

Probability of Recession

William Chiu

2022-11-25

Summary

Forecast the probability of a recession in the next 126 trading days using the following predictors:

1. Spread between 10Y CMT and Effective Federal Funds Rate
2. Lags of the spread
3. Adstock transformations of the spread

There are between 250 and 253 trading days in a year.

Extract Historical Data

Refer to this vignette for FRED data access.

```
library(tidyverse)
library(lubridate)
library(fredr)
library(car)
library(MLmetrics)
library(caret)
library(pdp)
library(gridExtra)
library(mboost)
library(gbm)
library(randomForest)
library(glmnet)
library(gtsummary)

randSeed <- 1983

startTestDate <- "1975-01-01"

# series_id <- c("FEDFUNDS", "GS10", "USREC", "UNRATE", "CPIAUCSL")

series_id <- c("DFF", "DGS10") # daily

response_id <- "USREC" # monthly

full_data <- map_dfr(series_id, function(x) {
  fredr(
```

```

    series_id = x,
    observation_start = as.Date("1950-01-01"),
    observation_end = as.Date("2022-12-01")
  )
})

recession_dates <- map_dfr(response_id, function(x) {
  fredr(
    series_id = x,
    observation_start = as.Date("1950-01-01"),
    observation_end = as.Date("2022-12-01")
  )
})

```

Pivot Wider

```

full_data_wide_raw <- full_data %>%
  arrange(date) %>%
  select(date, series_id, value) %>%
  pivot_wider(id_cols=date, names_from = series_id,
              values_from = value)%>%
  drop_na()

```

Calculate Features/Predictors

```

full_data_wide_features <- full_data_wide_raw %>%
  arrange(date) %>%
  mutate(SPRD_10YCMT_FEDFUNDS = DGS10 - DFF
         ) %>%
  mutate(across(
    .cols=c(SPRD_10YCMT_FEDFUNDS),
    .fns=list(lag1m = ~lag(.x, 1*21),
              lag3m = ~lag(.x, 3*21),
              lag6m = ~lag(.x, 6*21),
              lag9m = ~lag(.x, 9*21),
              lag12m = ~lag(.x, 12*21),
              lag5d = ~lag(.x, 5),
              lag10d = ~lag(.x, 10),
              lag15d = ~lag(.x, 15)
            )
  )) %>%
  drop_na()

```

Calculate Adstock

The adstock transformation is an auto-regressive transformation of a time series. The transformation takes into account past values of the time series. The intuition is that past values of the time series has a contemporaneous effect on the outcome.

$$AdStock(x_t) = x_t + \theta AdStock(x_{t-1})$$

where

$$0 < \theta < 1$$

The parameters cannot be estimated easily with least squares or logistic regression. Instead, we assume a range of potential values.

```
full_data_wide_features_adstock <- full_data_wide_features %>%
  arrange(date) %>%
  mutate(across(
    .cols=c(SPRD_10YCMT_FEDFUNDS),
    .fns=list(adstkL = ~stats::filter(.x,
                                          filter=0.01,
                                          method="recursive") ,
              adstkM = ~stats::filter(.x,
                                          filter=0.05,
                                          method="recursive") ,
              adstkHL = ~stats::filter(.x,
                                          filter=0.10,
                                          method="recursive"),
              adstkHM = ~stats::filter(.x,
                                          filter=0.15,
                                          method="recursive"),
              adstkHH = ~stats::filter(.x,
                                          filter=0.25,
                                          method="recursive"),
              adstkHHH = ~stats::filter(.x,
                                          filter=0.35,
                                          method="recursive")
    ))) %>%
  mutate(constant=1)
```

Calculate Moving Average

```
ma_fun <- function(k_param){
  rep(1/k_param, k_param)
}

full_data_wide_features_adstock <- full_data_wide_features_adstock %>%
  arrange(date) %>%
  mutate(across(
    .cols=c(SPRD_10YCMT_FEDFUNDS),
    .fns=list(ma5d = ~stats::filter(.x,
                                      filter=ma_fun(5),
                                      method="convolution",
                                      sides=1) ,
              ma10d = ~stats::filter(.x,
                                      filter=ma_fun(10),
                                      method="convolution",
```

```

        sides=1) ,
ma15d = ~stats::filter(.x,
                        filter=ma_fun(15),
                        method="convolution",
                        sides=1),
ma20d = ~stats::filter(.x,
                        filter=ma_fun(20),
                        method="convolution",
                        sides=1),
ma25d = ~stats::filter(.x,
                        filter=ma_fun(25),
                        method="convolution",
                        sides=1),
ma2m = ~stats::filter(.x,
                      filter=ma_fun(2*21),
                      method="convolution",
                      sides=1),
ma3m = ~stats::filter(.x,
                      filter=ma_fun(3*21),
                      method="convolution",
                      sides=1),
ma6m = ~stats::filter(.x,
                      filter=ma_fun(6*21),
                      method="convolution",
                      sides=1),
ma9m = ~stats::filter(.x,
                      filter=ma_fun(9*21),
                      method="convolution",
                      sides=1),
ma12m = ~stats::filter(.x,
                      filter=ma_fun(12*21),
                      method="convolution",
                      sides=1)
)))

```

Recession in next 6 months

```

full_data_wide <- full_data_wide_features_adstock %>%
  arrange(date) %>%
  mutate(date_month = month(date),
         date_year = year(date))

recession_df <- recession_dates %>%
  select(date, value) %>%
  arrange(date) %>%
  mutate(date_month = month(date),
         date_year = year(date))

full_data_wide <- full_data_wide %>%
  left_join(recession_df,

```

```

      by = c("date_month" = "date_month",
            "date_year" = "date_year")) %>%
mutate(USREC = value)

df_FUTREC = as.data.frame(
  data.table::shift(
    full_data_wide$USREC,
    n = 1:(6 * 21),
    type = "lead",
    give.names = TRUE,
    fill = NA
  )
) %>%
rowwise() %>%
mutate(FUTREC = max(c_across(V1_lead_1:V1_lead_126)))

full_data_wide$FUTREC <- df_FUTREC$FUTREC

full_data_wide <- full_data_wide %>%
  select(date=date.x, everything(), -date_month,
        -date_year, -date.y,
        -value) %>%
  drop_na()

full_data_wide_noUSREC <- full_data_wide %>%
  select(-USREC)

```

Remove the last 12 months of historical data

Since the NBER often dates recessions after they have already occurred (and sometimes ended), remove the last 12 months of historical data from both the training and test data sets.

```

recent_data <- tail(full_data_wide_noUSREC, 12*21)

train_test <- head(full_data_wide_noUSREC, -12*21)

```

Split Train/Test

```

full_size <- nrow(train_test)

test_size <- floor(full_size*0.5)

test_id <- seq.int(1,test_size,1)

train_test$constant <- 1

train_data <- train_test[-test_id,]
test_data <- train_test[test_id,]

```

```

train_yes_no <- train_data %>%
  mutate(FUTREC = case_when(FUTREC == 1 ~ "yes",
                             TRUE ~ "no"))

train_yes_no$FUTREC <- factor(train_yes_no$FUTREC,
                              levels=c("yes", "no"))

tbl_summary(train_data)

```

Characteristic	N = 7,157
date	1992-09-18 to 2021-04-28
DFF	1.91 (0.18, 4.92)
DGS10	4.06 (2.47, 5.49)
SPRD_10YCMT_FEDFUNDS	1.51 (0.58, 2.58)
SPRD_10YCMT_FEDFUNDS_lag1m	1.52 (0.58, 2.59)
SPRD_10YCMT_FEDFUNDS_lag3m	1.53 (0.58, 2.61)
SPRD_10YCMT_FEDFUNDS_lag6m	1.57 (0.58, 2.65)
SPRD_10YCMT_FEDFUNDS_lag9m	1.61 (0.59, 2.67)
SPRD_10YCMT_FEDFUNDS_lag12m	1.66 (0.61, 2.68)
SPRD_10YCMT_FEDFUNDS_lag5d	1.52 (0.58, 2.58)
SPRD_10YCMT_FEDFUNDS_lag10d	1.52 (0.58, 2.58)
SPRD_10YCMT_FEDFUNDS_lag15d	1.52 (0.58, 2.58)
SPRD_10YCMT_FEDFUNDS_adstkL	1.53 (0.59, 2.61)
SPRD_10YCMT_FEDFUNDS_adstkM	1.59 (0.61, 2.71)
SPRD_10YCMT_FEDFUNDS_adstkHL	1.68 (0.64, 2.86)
SPRD_10YCMT_FEDFUNDS_adstkHM	1.78 (0.68, 3.03)
SPRD_10YCMT_FEDFUNDS_adstkHH	2.02 (0.77, 3.44)
SPRD_10YCMT_FEDFUNDS_adstkHHH	2.32 (0.90, 3.97)
constant	7,157 (100%)
SPRD_10YCMT_FEDFUNDS_ma5d	1.51 (0.58, 2.58)
SPRD_10YCMT_FEDFUNDS_ma10d	1.51 (0.59, 2.59)
SPRD_10YCMT_FEDFUNDS_ma15d	1.50 (0.59, 2.59)
SPRD_10YCMT_FEDFUNDS_ma20d	1.50 (0.59, 2.60)
SPRD_10YCMT_FEDFUNDS_ma25d	1.50 (0.59, 2.61)
SPRD_10YCMT_FEDFUNDS_ma2m	1.51 (0.58, 2.61)
SPRD_10YCMT_FEDFUNDS_ma3m	1.51 (0.56, 2.62)
SPRD_10YCMT_FEDFUNDS_ma6m	1.55 (0.57, 2.62)
SPRD_10YCMT_FEDFUNDS_ma9m	1.56 (0.55, 2.66)
SPRD_10YCMT_FEDFUNDS_ma12m	1.57 (0.55, 2.69)
FUTREC	959 (13%)

Remove stale data from test set

Exclude historical data prior to 1975-01-01 because the economy changed dramatically (due to computational innovation).

```
summary(test_data$date)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
----	------	---------	--------	------	---------	------

```
## "1964-01-09" "1971-03-14" "1978-05-16" "1978-05-16" "1985-07-23" "1992-09-17"
```

```
test_data <- test_data %>%  
  filter(date >= startTestDate)  
  
summary(test_data$date)
```

```
##           Min.         1st Qu.         Median         Mean         3rd Qu.         Max.  
## "1975-01-02" "1979-06-09" "1983-11-14" "1983-11-13" "1988-04-20" "1992-09-17"
```

Setup Parallel Processing

```
library(doParallel)  
  
cl <- makePSOCKcluster(4)  
registerDoParallel(cl)
```

Cross-Validation Framework

```
foldes <- 20  
fcstHorizon <- 12*21  
initWindow <- nrow(train_yes_no) - fcstHorizon * foldes  
  
if(initWindow < 100){  
  stop("Too few observations.")  
}  
  
fitControl_oneSE <- trainControl(method = "timeslice",  
                                initialWindow=initWindow,  
                                horizon=fcstHorizon,  
                                fixedWindow=FALSE,  
                                skip=fcstHorizon-1,  
                                ## Estimate class probabilities  
                                classProbs = TRUE,  
                                ## Evaluate performance using  
                                ## the following function  
                                summaryFunction = mnLogLoss,  
                                selectionFunction="oneSE")  
  
fitControl_best <- trainControl(method = "timeslice",  
                                initialWindow=initWindow,  
                                horizon=fcstHorizon,  
                                fixedWindow=FALSE,  
                                skip=fcstHorizon-1,  
                                ## Estimate class probabilities  
                                classProbs = TRUE,  
                                ## Evaluate performance using  
                                ## the following function  
                                summaryFunction = mnLogLoss,  
                                selectionFunction="best")
```


Gradient Boosting for Additive Models

```
grid_gam <- expand.grid(mstop=seq(1,50,2),
                       prune="no")

set.seed(randSeed)

gam_mod <- train(
  FUTREC ~ . - date - constant,
  data = train_yes_no,
  method = "gamboost",
  trControl = fitControl_best,
  metric = "logLoss",
  tuneGrid = grid_gam,
  family = Binomial()
)

plot(gam_mod)
```



```
gam_mod$bestTune
```

```
##  mstop prune
##  9      17  no
```



eXtreme Gradient Boosting Trees

```
grid_xgb <- expand.grid(nrounds=seq(1, 20, 2),
  max_depth=c(1,3,5),
  eta=seq(0.05,1,0.1),
  gamma=0,
  colsample_bytree=1,
  min_child_weight=10,
  subsample=1
)
```

```
set.seed(randSeed)
```

```
xgb_mod <- train(
  FUTREC ~ . - date - constant,
  data = train_yes_no,
  method = "xgbTree",
  trControl = fitControl_best,
  metric = "logLoss",
  tuneGrid = grid_xgb,
  objective = "binary:logistic"
)
```

```
plot(xgb_mod)
```



```
xgb_mod$bestTune
```

```
##      nrounds max_depth eta gamma colsample_bytree min_child_weight subsample
## 37         13         1 0.15         0             1             10         1
```


Random Forest

```
grid_rf <- data.frame(mtry=seq.int(1,50,5))

set.seed(randSeed)

rf_mod <- train(
  FUTREC ~ . - date - constant,
  data = train_yes_no,
  method = "rf",
  trControl = fitControl_best,
  metric = "logLoss",
  tuneGrid = grid_rf,
  importance = TRUE
)

plot(rf_mod)
```



```
rf_mod$bestTune
```

```
##      mtry  
## 1      1
```

Stepwise Regression

The `glmStepAIC` method uses the `glm()` function from the `stats` package. The documentation for `glm()` says:

For binomial and quasibinomial families the response can also be specified as a factor (when the first level denotes failure and all others success) or as a two-column matrix with the columns giving the numbers of successes and failures.

However, for most methods (that do not invoke `glm()`) in `train`, the first level denotes the success (the opposite of `glm()`). This behavior causes the coefficient signs to flip. Be highly suspicious when interpreting coefficients from models that are fit using `train`.

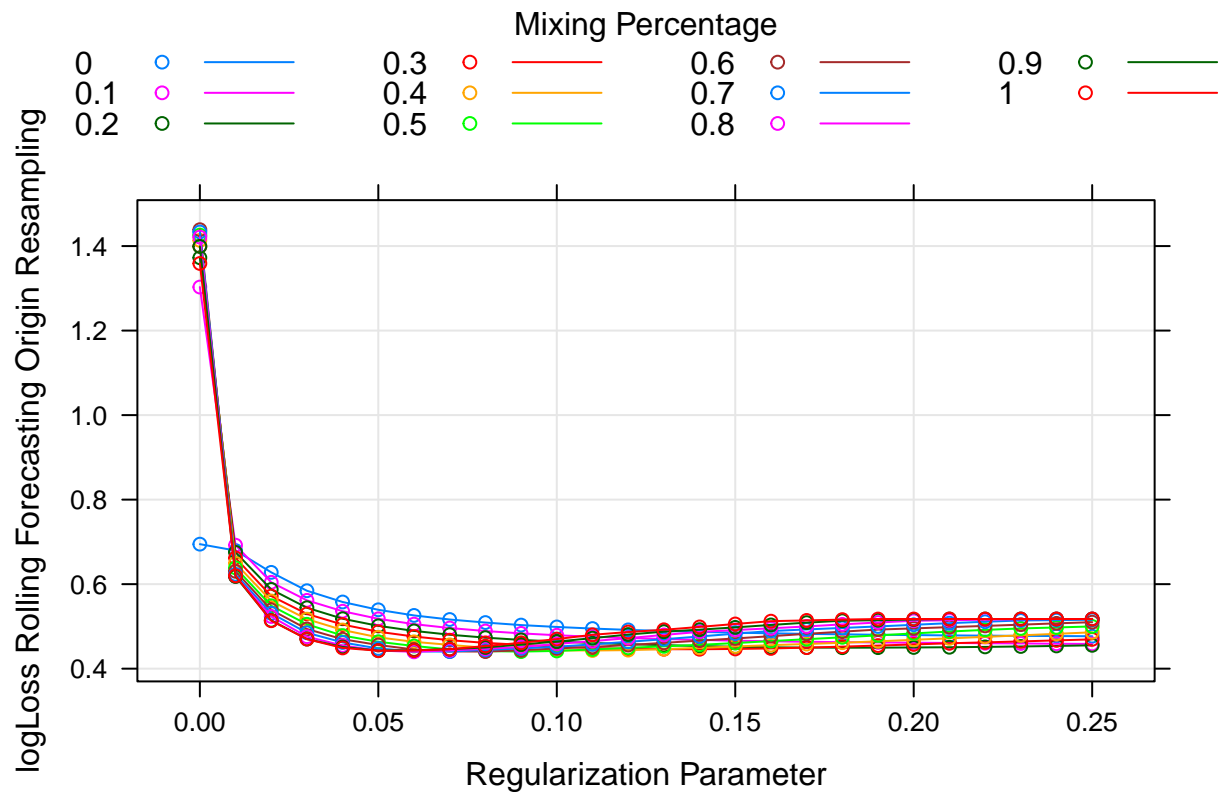
```
set.seed(randSeed)  
  
stepwise_mod <- train(  
  FUTREC ~ . - date - constant,  
  data = train_yes_no,  
  method = "glmStepAIC",  
  trControl = fitControl_best,  
  metric = "logLoss",  
  tuneLength = 10,  
  family = binomial,  
  trace = 0,  
  k = 10*log(nrow(train_yes_no)),  
  direction = "forward"  
)
```

Elastic Net (Lasso)

```
grid_glmnet <- expand.grid(  
  alpha = seq(0,1,0.1),  
  lambda = seq(0, 0.25, 0.01)  
)  
  
set.seed(randSeed)  
  
glmnet_mod <- train(  
  FUTREC ~ . - date - constant,  
  data = train_yes_no,  
  method = "glmnet",  
  trControl = fitControl_best,  
  metric = "logLoss",  
  tuneGrid = grid_glmnet,
```

```
family = "binomial"
)

plot(glmnet_mod)
```



```
glmnet_mod$bestTune
```

```
##      alpha lambda
## 215    0.8    0.06
```

Multivariate Adaptive Regression Splines

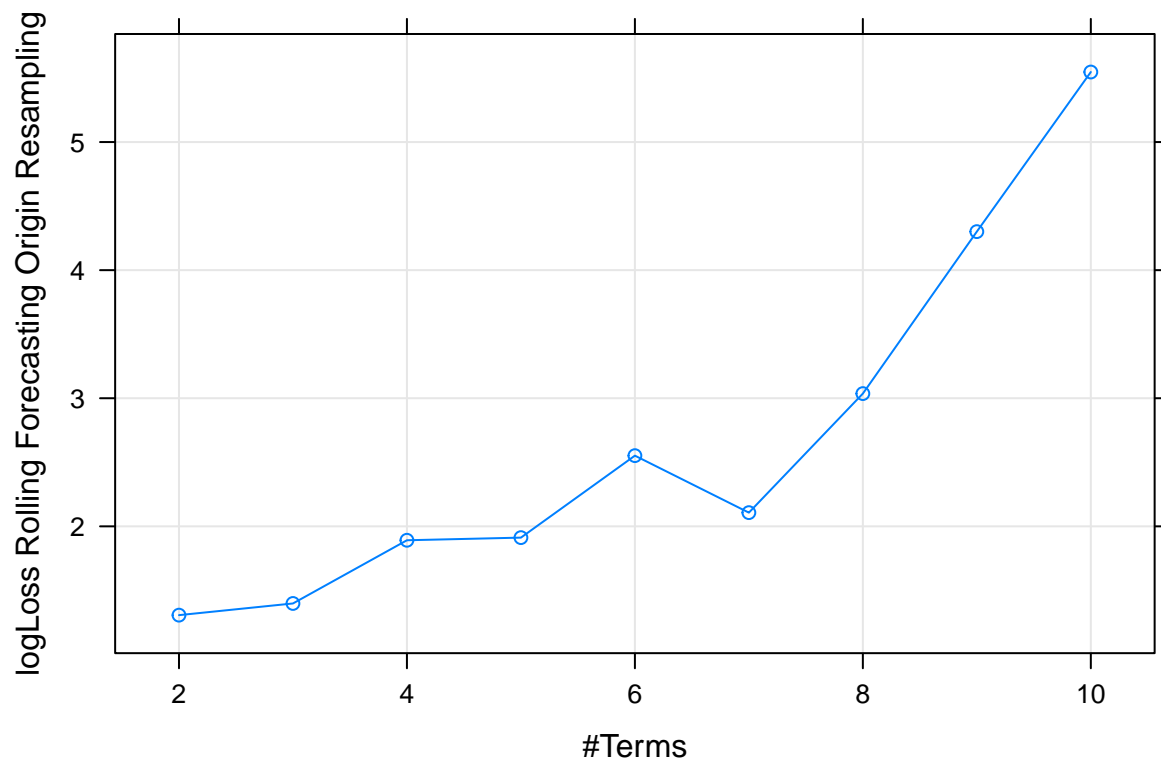
```
grid_mars <- expand.grid(nprune=seq(2,10,1),
                        degree=1)

set.seed(randSeed)

earth_mod <- train(
  FUTREC ~ . - date - constant,
  data = train_yes_no,
  method = "earth",
  trControl = fitControl_best,
```

```
metric = "logLoss",
tuneGrid = grid_mars,
glm = list(family = binomial)
)

plot(earth_mod)
```



```
earth_mod$bestTune
```

```
##  nprune degree
## 1      2      1
```

Null Model: Intercept-only Model

```
set.seed(randSeed)

null_mod <- train(
  FUTREC ~ constant,
  data = train_yes_no,
  method = "glm",
  trControl = fitControl_best,
  metric = "logLoss",
```


```
family = binomial
)
```

Compare Models

```
resamps <- resamples(list(XGB = xgb_mod,
                          GAM = gam_mod,
                          RF = rf_mod,
                          Step = stepwise_mod,
                          Lasso = glmnet_mod,
                          MARS = earth_mod,
                          Null = null_mod)
)
summary(resamps)
```

```
##
## Call:
## summary.resamples(object = resamps)
##
## Models: XGB, GAM, RF, Step, Lasso, MARS, Null
## Number of resamples: 20
##
## logLoss
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## XGB  9.192479e-02 0.11526830 0.1304653 0.3986430 0.4869654 1.703535    0
## GAM  4.987642e-02 0.08323740 0.1231974 0.4190890 0.4005984 2.122464    0
## RF   7.040866e-03 0.06073099 0.2466586 0.9677976 0.6102503 10.913992    0
## Step 2.022540e-07 0.05006111 0.2384969 1.5420435 0.5442494 17.817623    0
## Lasso 1.422861e-02 0.07611252 0.1499059 0.4398014 0.4424810 2.242503    0
## MARS  7.981488e-05 0.05497337 0.0712627 1.3066759 0.6221208 17.817623    0
## Null  9.011199e-02 0.13169314 0.1611508 0.5175254 0.2956142 2.193465    0
```

```
dotplot(resamps, metric = "logLoss", conf.level=0.25)
```



Explore XGB Model

```
xgb_mod$bestTune
```

```
##      nrounds max_depth  eta gamma colsample_bytree min_child_weight subsample
## 37         13         1 0.15    0                   1                10         1
```

```
df_imp <- varImp(xgb_mod)$importance %>%
  arrange(desc(Overall))
```

```
df_imp$variable <- rownames(df_imp)
```

```
df_imp <- df_imp %>%
  select(variable, Overall)
```

```
row.names(df_imp) <- NULL
```

```
knitr::kable(df_imp)
```

variable	Overall
SPRD_10YCMT_FEDFUNDS_lag6m	100.00000

variable	Overall
SPRD_10YCMT_FEDFUNDS_lag12m	79.69336
SPRD_10YCMT_FEDFUNDS_lag9m	63.85953
SPRD_10YCMT_FEDFUNDS_ma9m	53.70159
DFD	0.00000
DGS10	0.00000
SPRD_10YCMT_FEDFUNDS	0.00000
SPRD_10YCMT_FEDFUNDS_lag1m	0.00000
SPRD_10YCMT_FEDFUNDS_lag3m	0.00000
SPRD_10YCMT_FEDFUNDS_lag5d	0.00000
SPRD_10YCMT_FEDFUNDS_lag10d	0.00000
SPRD_10YCMT_FEDFUNDS_lag15d	0.00000
SPRD_10YCMT_FEDFUNDS_adstkL	0.00000
SPRD_10YCMT_FEDFUNDS_adstkM	0.00000
SPRD_10YCMT_FEDFUNDS_adstkHL	0.00000
SPRD_10YCMT_FEDFUNDS_adstkHM	0.00000
SPRD_10YCMT_FEDFUNDS_adstkHH	0.00000
SPRD_10YCMT_FEDFUNDS_adstkHHH	0.00000
SPRD_10YCMT_FEDFUNDS_ma5d	0.00000
SPRD_10YCMT_FEDFUNDS_ma10d	0.00000
SPRD_10YCMT_FEDFUNDS_ma15d	0.00000
SPRD_10YCMT_FEDFUNDS_ma20d	0.00000
SPRD_10YCMT_FEDFUNDS_ma25d	0.00000
SPRD_10YCMT_FEDFUNDS_ma2m	0.00000
SPRD_10YCMT_FEDFUNDS_ma3m	0.00000
SPRD_10YCMT_FEDFUNDS_ma6m	0.00000
SPRD_10YCMT_FEDFUNDS_ma12m	0.00000

```

pdp.top1 <- partial(xgb_mod,
  pred.var = df_imp$variable[1],
  plot = TRUE,
  rug = TRUE)

pdp.top2 <- partial(xgb_mod,
  pred.var = df_imp$variable[2],
  plot = TRUE,
  rug = TRUE)

pdp.top3 <- partial(xgb_mod,
  pred.var = df_imp$variable[3],
  plot = TRUE,
  chull = TRUE
)

pdp.top4 <- partial(xgb_mod,
  pred.var = df_imp$variable[4],
  plot = TRUE,
  chull = TRUE
)

pdp.top5 <- partial(xgb_mod,
  pred.var = df_imp$variable[5],

```



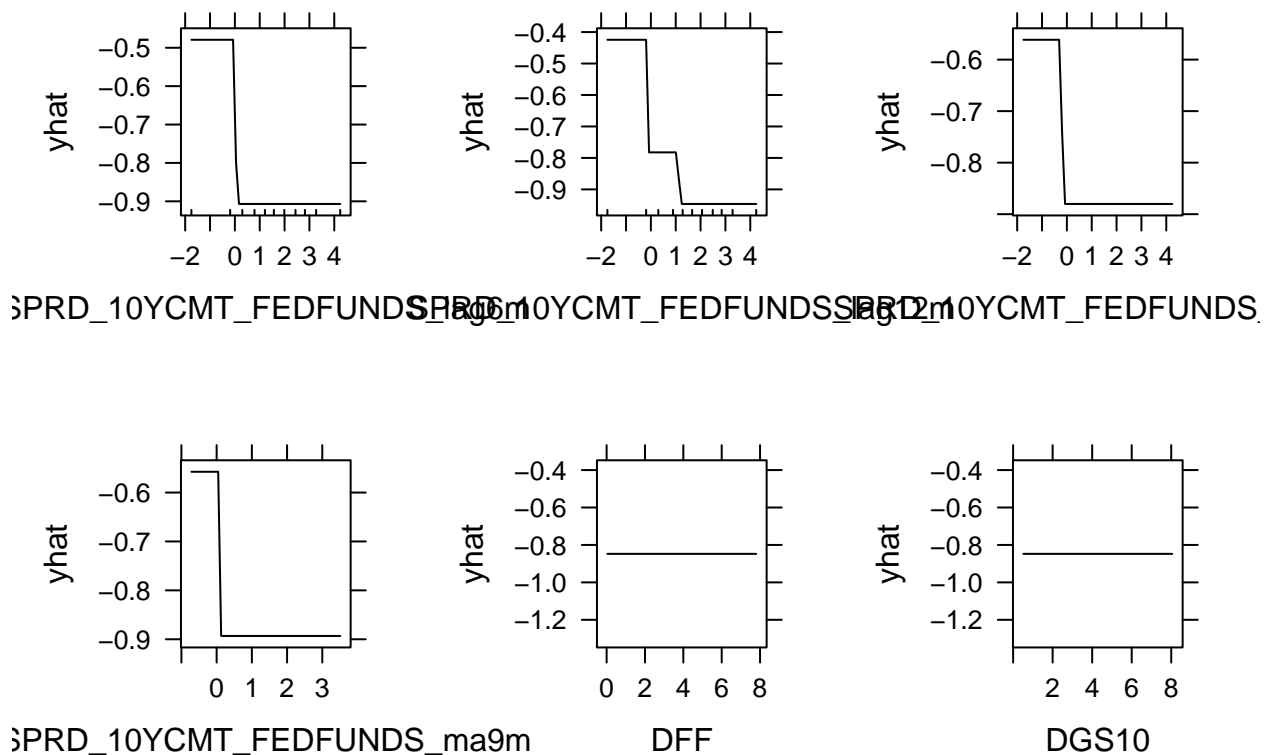
```

    plot = TRUE,
    chull = TRUE
  )

pdp.top6 <- partial(xgb_mod,
  pred.var = df_imp$variable[6],
  plot = TRUE,
  chull = TRUE
)

grid.arrange(pdp.top1, pdp.top2, pdp.top3,
  pdp.top4, pdp.top5, pdp.top6, ncol = 3)

```



Peeking

Peeking means we use the insights from the automated models to choose variables in subsequent models. This is technically cheating and causes the cross-validation errors to be artificially low. This is addressed in the test set which does not have peeking bias.

```

top_predictors <- head(df_imp$variable)

best_predictor <- head(top_predictors, 1)

top_fm1a <- as.formula(paste0("FUTREC ~",

```

```

                                paste0(top_predictors,
                                           collapse=" + "))

top1_fm1a <- as.formula(paste0("FUTREC ~",
                                paste0(best_predictor,
                                           collapse=" + "))

```

Logistic Regression (with peeking)

As mentioned early, `train` and `glm` treat the reference level differently for binary outcomes. Hence, the coefficients are flipped when training a logistic regression inside `train`.

```

logit_mod <- train(
  top1_fm1a,
  data = train_yes_no,
  method = "glm",
  trControl = fitControl_best,
  metric = "logLoss",
  family=binomial
)

summary(logit_mod)

```

```

##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.3929   0.1184   0.2504   0.4903   1.9110
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      0.64204    0.04400   14.59  <2e-16 ***
## SPRD_10YCMT_FEDFUNDS_lag6m  1.34846    0.04443   30.35  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 5638.4  on 7156  degrees of freedom
## Residual deviance: 4144.0  on 7155  degrees of freedom
## AIC: 4148
##
## Number of Fisher Scoring iterations: 6

```

Compare Models

CV errors for models with peeking are misleadingly low. This will be addressed with a test set.

```

mymods <- list(XGB = xgb_mod,
               GAM = gam_mod,
               RF = rf_mod,
               Step = stepwise_mod,
               Lasso = glmnet_mod,
               MARS = earth_mod,
               Null = null_mod,
               Logit = logit_mod) ## peeking

resamps <- resamples(mymods)
summary(resamps)

```

```


##
## Call:
## summary.resamples(object = resamps)
##
## Models: XGB, GAM, RF, Step, Lasso, MARS, Null, Logit
## Number of resamples: 20
##
## logLoss
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## XGB  9.192479e-02 0.11526830 0.13046527 0.3986430 0.4869654  1.703535    0
## GAM  4.987642e-02 0.08323740 0.12319739 0.4190890 0.4005984  2.122464    0
## RF   7.040866e-03 0.06073099 0.24665857 0.9677976 0.6102503 10.913992    0
## Step 2.022540e-07 0.05006111 0.23849694 1.5420435 0.5442494 17.817623    0
## Lasso 1.422861e-02 0.07611252 0.14990590 0.4398014 0.4424810  2.242503    0
## MARS  7.981488e-05 0.05497337 0.07126270 1.3066759 0.6221208 17.817623    0
## Null  9.011199e-02 0.13169314 0.16115082 0.5175254 0.2956142  2.193465    0
## Logit 5.135720e-05 0.03049939 0.09049806 0.4547867 0.4109226  4.244014    0

```

```

dotplot(resamps, metric = "logLoss", conf.level=0.25)

```



Test Set Performance

```
perf <-
  function(lst_mods,
           f_metric = caTools::colAUC,
           metricname = "ROC-AUC",
           dat=test_data,
           response="FUTREC") {
    lst_preds <- map(
      .x = lst_mods,
      .f = function(x) {
        if (class(x)[1] != "train") {
          predict(x, newdata = dat, type = "response")
        } else
          (
            predict(x, newdata = dat, type = "prob")[, "yes"]
          )
      }
    )
    map_dfr(lst_preds, function(x) {
      f_metric(x, dat[,response, drop=TRUE])
    }) %>%
```

```

    pivot_longer(everything(), names_to = "model", values_to = metricname)
  }

perf(mymods, caTools::colAUC, "ROC-AUC") %>%
  arrange(desc(`ROC-AUC`)) %>%
  knitr::kable()

```

model	ROC-AUC
XGB	0.9308781
GAM	0.9275417
Lasso	0.9206350
Logit	0.9058277
MARS	0.8478348
Step	0.8170368
RF	0.7931440
Null	0.5000000

```

perf(mymods, MLmetrics::LogLoss, "LogLoss") %>%
  arrange(LogLoss) %>%
  knitr::kable()

```

model	LogLoss
XGB	0.3282365
Lasso	0.3619716
Logit	0.3764121
GAM	0.4657387
RF	0.5407493
Null	0.5885335
Step	0.7092285
MARS	1.0010820

Probability of Recession (Most Recent Trading Day)

```

curr_data <- tail(full_data_wide_features_adstock, 1)

curr_data$date

```

```
## [1] "2022-11-23"
```

```

score_fun <- function(mods, dat) {
  output <- map_dfc(.x = mods, .f = function(x) {
    if(class(x)[1] != "train"){
      predict(x, newdata = dat, type = "response")
    } else(
      predict(x, newdata = dat, type = "prob")[,"yes"]
    )
  })
}

```

```

}) %>%
  pivot_longer(everything(), names_to = "model",
               values_to = "prob_rec")

output$prob_rec <- scales::percent(output$prob_rec)

return(output)
}

knitr::kable(score_fun(mymods, curr_data))

```

model	prob_rec
XGB	10.27%
GAM	7.57%
RF	1.20%
Step	1.61%
Lasso	8.66%
MARS	5.46%
Null	13.40%
Logit	3.29%

Backtesting

```

full_data_bktst <- full_data_wide %>%
  filter(date >= startTestDate)

bkst_fun <- function(mods, dat) {
  output <- map_dfc(.x = mods, .f = function(x) {
    if(class(x)[1] != "train"){
      predict(x, newdata = dat, type = "response")
    } else(
      predict(x, newdata = dat, type = "prob")[, "yes"]
    )
  })

  output$date <- dat$date

  output <- output%>%
    pivot_longer(-date, names_to = "model",
                 values_to = "prob_rec")

  return(output)
}

df_plot <- bkst_fun(mymods, full_data_bktst)

```

```

actuals <- full_data_bktst %>%
  mutate(model="actuals") %>%
  select(date, model, prob_rec=USREC)

df_plot_final <- bind_rows(df_plot, actuals)

end_test_date <- max(test_data$date)

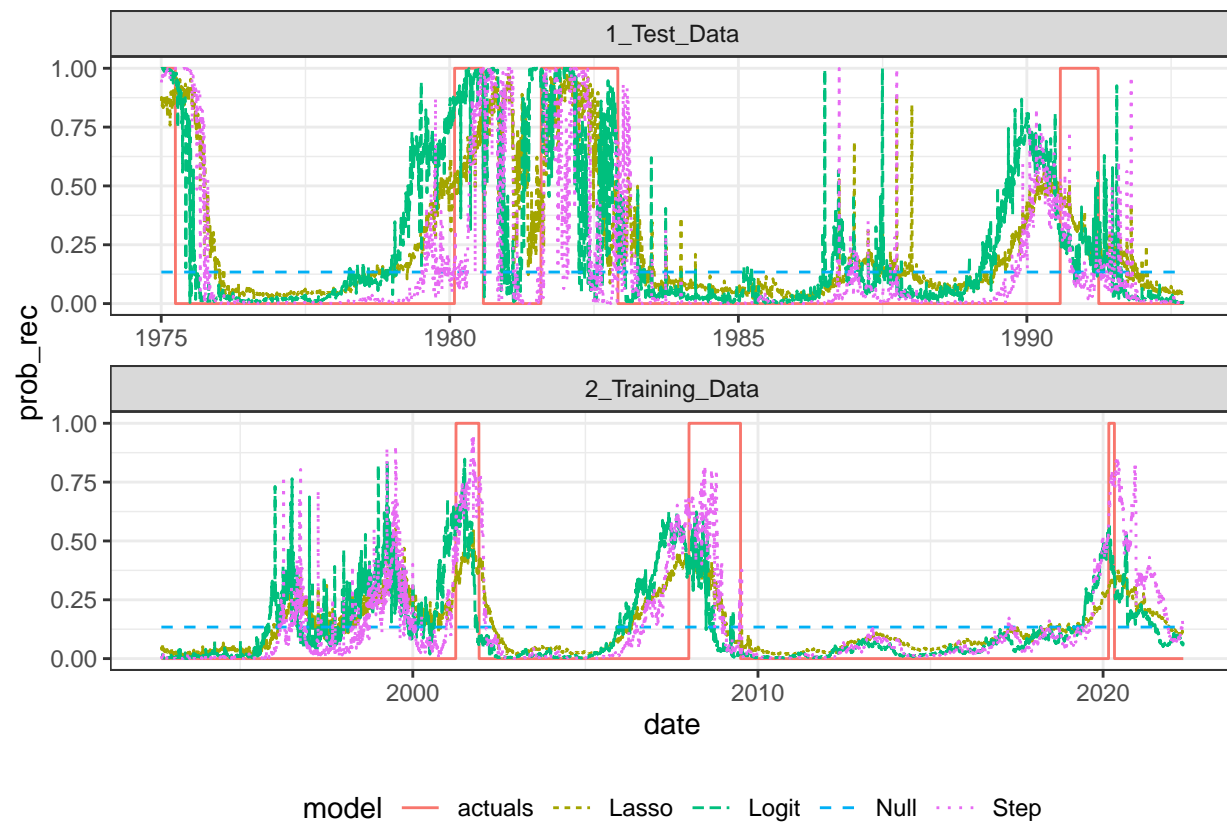
df_plot_final <- df_plot_final %>%
  mutate(epoc = case_when(date <= end_test_date ~ "1_Test_Data",
                           TRUE ~ "2_Training_Data")
  )

df_plot_logit_scam <- df_plot_final %>%
  filter(model %in% c('actuals', 'Null',
                     'Logit', 'Step', 'Lasso',
                     'LogitKnot'))

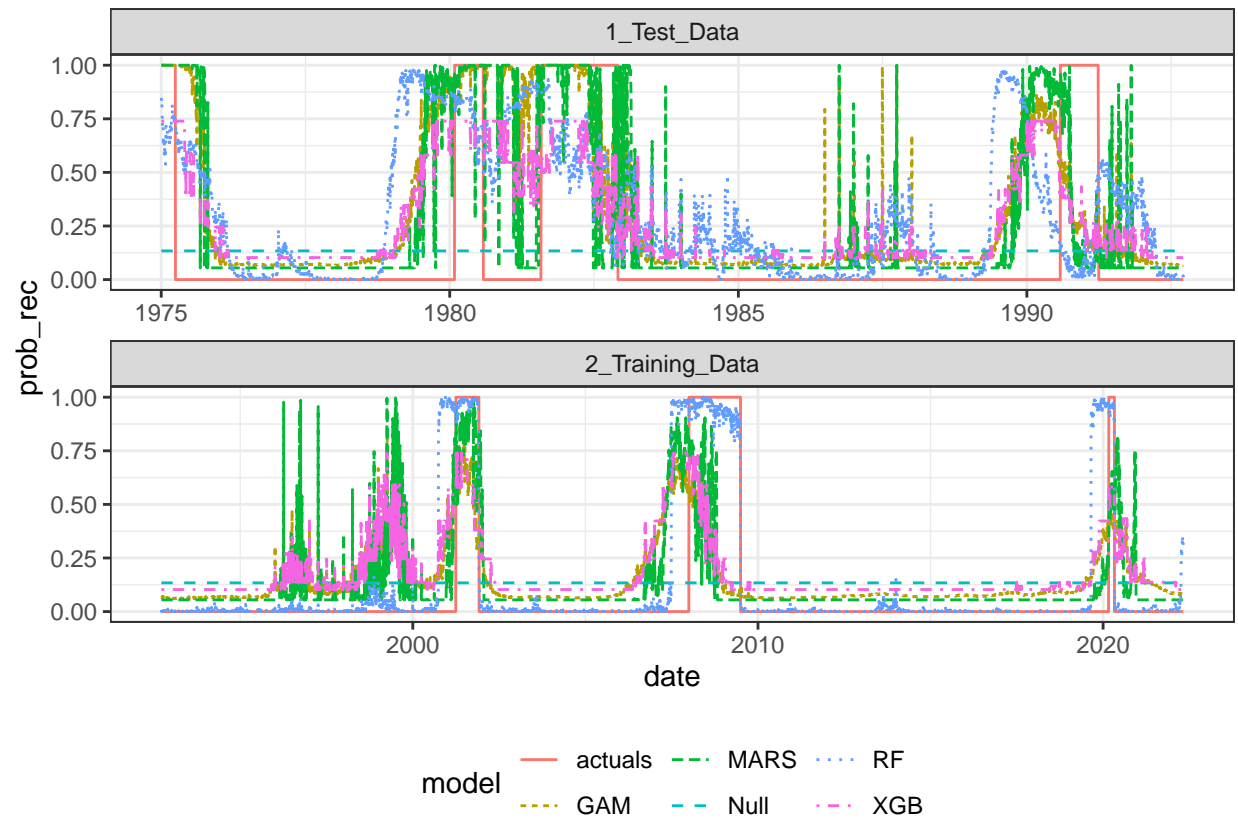
df_plot_knots_gbm <- df_plot_final %>%
  filter(model %in% c('actuals', 'Null',
                     'XGB', 'RF',
                     'GAM',
                     'MARS'))

ggplot(df_plot_logit_scam, aes(x=date, y=prob_rec, group=model,
                              linetype=model, color=model)) +
  geom_line() +
  theme_bw() +
  theme(legend.position = "bottom") +
  facet_wrap(vars(epoc), scales="free", nrow=2)

```



```
ggplot(df_plot_knots_gbm, aes(x=date, y=prob_rec, group=model,
                             linetype=model, color=model)) +
  geom_line() +
  theme_bw() +
  theme(legend.position = "bottom") +
  facet_wrap(vars(epoc), scales="free", nrow=2)
```

```
stopCluster(c1)
```